

MLOps Pipeline with Kubeflow and MLflow:

A Comprehensive Guide to Production-Ready ML Infrastructure

Vladimir Ovcharov

MLOps Engineer & ML Systems Architect

vladimir.ovcharov@highfunk.uk

July 10, 2025

Abstract

This article presents a comprehensive guide to building production-ready machine learning pipelines using Kubeflow for orchestration and MLflow for experiment tracking. We explore the architectural patterns, implementation strategies, and best practices for creating scalable, maintainable, and robust MLOps infrastructure. The guide includes practical code examples, deployment configurations, and monitoring strategies that can be directly applied in enterprise environments.

Keywords: MLOps, Kubeflow, MLflow, Kubernetes, Machine Learning Pipeline, Experiment Tracking, Model Deployment

1 Introduction

The evolution of machine learning from research prototypes to production-ready systems represents one of the most significant challenges facing modern data science organizations. While the academic community has made tremendous strides in developing sophisticated algorithms and achieving state-of-the-art results on benchmark datasets, the transition to real-world applications introduces a complex web of operational, technical, and organizational challenges that traditional software engineering practices struggle to address effectively.

1.1 The Production Challenge

The journey from a Jupyter notebook experiment to a production machine learning system involves far more than simply deploying a trained model. Organizations typically encounter what has been termed the "ML production gap" – the substantial difference between the controlled environment of model development and the chaotic reality of production systems where models must operate reliably, scale dynamically, and maintain performance over time.

Consider the typical lifecycle of a machine learning project: data scientists begin with exploratory data analysis, experiment with various algorithms and feature engineering approaches, tune hyperparameters, and eventually arrive at a model that demonstrates promising performance on held-out test data. However, this process often occurs in isolation, using static datasets, controlled environments, and manual workflows that bear little resemblance to the dynamic, distributed, and automated systems required for production deployment.

The challenges multiply when we consider the operational requirements of production ML systems:

- **Data Pipeline Reliability:** Production systems must handle streaming data, data quality issues, schema evolution, and upstream system failures gracefully

- **Model Performance Monitoring:** Unlike traditional software where bugs are typically deterministic, ML models can degrade silently due to data drift, concept drift, or changing business conditions
- **Scalability Requirements:** Models must serve predictions at scale, often handling thousands of requests per second with strict latency requirements
- **Regulatory Compliance:** Many industries require explainability, auditability, and bias detection capabilities that are difficult to retrofit into existing systems
- **Continuous Learning:** Models must be retrained and updated regularly to maintain performance, requiring sophisticated automation and validation pipelines

1.2 The Fragmentation Problem

Traditional approaches to ML deployment have led to significant fragmentation across the ML lifecycle. Different teams often use disparate tools for data preparation, model training, validation, deployment, and monitoring. This fragmentation creates several critical issues:

Tool Proliferation: Organizations frequently find themselves managing dozens of different tools and platforms, each optimized for a specific phase of the ML lifecycle. Data engineers might use Apache Airflow for data pipelines, data scientists prefer Jupyter notebooks for experimentation, ML engineers deploy models using custom Docker containers, and operations teams monitor systems using traditional APM tools that lack ML-specific metrics.

Knowledge Silos: The fragmented toolchain often leads to knowledge silos where each team becomes expert in their specific tools but lacks understanding of the broader system. This creates bottlenecks, communication barriers, and makes it difficult to optimize the end-to-end workflow.

Reproducibility Crisis: Without standardized environments and workflows, reproducing experimental results becomes extremely challenging. Models that perform well in development may fail in production due to subtle differences in data preprocessing, library versions, or infrastructure configurations.

Deployment Bottlenecks: The handoff between data science teams and production systems often becomes a significant bottleneck. Models developed in Python on local machines must be translated into production-ready services, often requiring significant engineering effort and introducing opportunities for errors.

1.3 The MLOps Revolution

Machine Learning Operations (MLOps) has emerged as a discipline specifically designed to address these challenges by applying DevOps principles and practices to machine learning workflows. MLOps recognizes that ML systems are fundamentally different from traditional software applications and require specialized approaches to testing, deployment, monitoring, and maintenance.

The core principles of MLOps include:

1. **Automation:** Minimizing manual interventions through automated pipelines for data processing, model training, validation, and deployment
2. **Reproducibility:** Ensuring that every experiment, training run, and deployment can be exactly reproduced through proper versioning of code, data, models, and infrastructure

3. **Monitoring:** Implementing comprehensive monitoring that goes beyond traditional system metrics to include model-specific indicators like prediction drift, data quality, and business impact
4. **Collaboration:** Breaking down silos between data science, engineering, and operations teams through shared tools, processes, and vocabulary
5. **Governance:** Establishing controls and processes for model approval, deployment, and compliance with regulatory requirements

1.4 Technology Landscape and Solution Selection

The MLOps ecosystem has rapidly evolved to include numerous platforms, frameworks, and tools, each addressing different aspects of the ML production challenge. The selection of appropriate technologies requires careful consideration of organizational needs, existing infrastructure, team capabilities, and long-term strategic goals.

Kubeflow has emerged as a leading platform for ML workflows on Kubernetes, providing a comprehensive suite of tools that span the entire ML lifecycle. Built on the foundation of Kubernetes, Kubeflow inherits the scalability, reliability, and ecosystem benefits of the container orchestration platform while adding ML-specific capabilities:

- **Kubeflow Pipelines:** A platform for building and deploying portable, scalable ML workflows
- **Katib:** Automated hyperparameter tuning and neural architecture search
- **KServe:** Model serving platform with advanced features like canary deployments and multi-framework support
- **Notebooks:** Managed Jupyter notebook environments with resource allocation and sharing capabilities
- **Training Operators:** Distributed training support for TensorFlow, PyTorch, and other frameworks

MLflow complements Kubeflow by providing robust experiment tracking and model management capabilities. Originally developed by Databricks, MLflow has become a de facto standard for ML lifecycle management, offering:

- **Experiment Tracking:** Comprehensive logging of parameters, metrics, and artifacts for every experiment
- **Model Registry:** Centralized model store with versioning, staging, and annotation capabilities
- **Model Packaging:** Standardized format for packaging models with their dependencies
- **Model Serving:** Simple deployment options for various serving platforms

1.5 Integration Strategy and Benefits

The combination of Kubeflow and MLflow creates a powerful MLOps platform that addresses the full spectrum of production ML challenges. This integration strategy leverages the strengths of both platforms while mitigating their individual limitations:

Kubeflow provides the infrastructure and orchestration layer, handling the complex task of managing distributed workloads, resource allocation, and pipeline execution across Kubernetes clusters. Its native integration with Kubernetes means that ML workloads can benefit from the same scalability, reliability, and operational practices used for other cloud-native applications.

MLflow serves as the metadata and lifecycle management layer, providing the tracking, versioning, and governance capabilities essential for maintaining reproducibility and compliance in production environments. Its framework-agnostic approach ensures that teams can continue using their preferred ML libraries while benefiting from standardized lifecycle management.

The synergy between these platforms enables several key capabilities:

- **End-to-End Traceability:** Every model deployed in production can be traced back to its training data, code version, hyperparameters, and experimental results
- **Automated Retraining:** Pipelines can automatically detect model performance degradation and trigger retraining workflows with minimal human intervention
- **A/B Testing and Gradual Rollouts:** New model versions can be safely deployed using canary deployments and traffic splitting capabilities
- **Resource Optimization:** Kubernetes-native resource management ensures efficient utilization of computational resources across training and serving workloads
- **Multi-Environment Consistency:** The same pipeline definitions can be deployed across development, staging, and production environments with environment-specific configurations

1.6 Article Scope and Objectives

This comprehensive guide provides practical, implementable solutions for organizations seeking to establish robust MLOps practices using Kubeflow and MLflow. Rather than focusing on theoretical concepts, we emphasize hands-on implementation with real-world code examples, configuration files, and architectural patterns that have been proven in production environments.

The article is structured to take readers through a complete implementation journey, from initial environment setup through advanced monitoring and optimization techniques. Each section builds upon previous concepts while providing sufficient detail for independent implementation. Code examples are production-ready and include error handling, logging, and best practices gleaned from real-world deployments.

Our target audience includes ML engineers, DevOps practitioners, data scientists, and technical leaders who are responsible for moving ML systems from experimentation to production. We assume familiarity with basic ML concepts, Kubernetes fundamentals, and Python programming, but provide sufficient context for readers to understand and adapt the solutions to their specific environments.

By the end of this guide, readers will have a complete understanding of how to:

- Design and implement scalable ML pipelines using Kubeflow
- Establish comprehensive experiment tracking and model management with MLflow
- Deploy models safely and efficiently using modern serving platforms
- Monitor ML systems for performance, drift, and operational issues

- Implement automated retraining and continuous deployment workflows
- Apply security, governance, and compliance best practices

2 Implementation Guide

This section provides detailed, step-by-step instructions for implementing a production-ready MLOps platform using Kubeflow and MLflow. The implementation follows a progressive approach, starting with foundational infrastructure and gradually building up to advanced features. Each step includes comprehensive code examples, configuration files, and troubleshooting guidance based on real-world deployment experiences.

2.1 Prerequisites and Environment Preparation

Before beginning the implementation, ensure that your environment meets the necessary requirements and that all prerequisite tools are properly configured.

2.1.1 Infrastructure Requirements

Kubernetes Cluster Specifications:

- **Minimum cluster size:** 3 nodes with 4 CPU cores and 16GB RAM each
- **Recommended cluster size:** 5+ nodes with 8 CPU cores and 32GB RAM each
- **Kubernetes version:** 1.24 or later (tested up to 1.28)
- **Storage:** Dynamic volume provisioning with SSD-backed storage classes
- **Network:** CNI-compatible networking (Calico, Flannel, or cloud provider CNI)
- **Load Balancer:** Cloud provider load balancer or MetalLB for on-premises

Additional Infrastructure Components:

- Object storage (AWS S3, Google Cloud Storage, Azure Blob, or MinIO)
- Container registry (Docker Hub, ECR, GCR, or Harbor)
- DNS management for custom domains and SSL certificates
- Monitoring infrastructure (Prometheus operator recommended)

2.1.2 Required Tools and Dependencies

Install and configure the following tools on your management workstation:

```
1 #!/bin/bash
2
3 # Install kubectl
4 curl -LO "https://dl.k8s.io/release/$(curl -L -s https://dl.k8s.io/release/
   stable.txt)/bin/linux/amd64/kubectl"
5 sudo install -o root -g root -m 0755 kubectl /usr/local/bin/kubectl
6
7 # Install Helm
8 curl https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3 | bash
9
```

```

10 # Install kustomize
11 curl -s "https://raw.githubusercontent.com/kubernetes-sigs/kustomize/master/
    hack/install_kustomize.sh" | bash
12 sudo mv kustomize /usr/local/bin/
13
14 # Install yq for YAML processing
15 sudo wget -qO /usr/local/bin/yq https://github.com/mikefarah/yq/releases/latest
    /download/yq_linux_amd64
16 sudo chmod +x /usr/local/bin/yq
17
18 # Install kubens and kubectl for easier cluster management
19 sudo git clone https://github.com/ahmetb/kubectx /opt/kubectx
20 sudo ln -s /opt/kubectx/kubectx /usr/local/bin/kubectx
21 sudo ln -s /opt/kubectx/kubens /usr/local/bin/kubens
22
23 # Verify installations
24 echo "Verifying tool installations..."
25 kubectl version --client
26 helm version
27 kustomize version
28 yq --version

```

Listing 1: Tool Installation Script

2.2 Kubeflow Installation and Configuration

Kubeflow installation requires careful attention to component selection and configuration to ensure optimal performance and security.

2.2.1 Namespace and Security Setup

Begin by creating the necessary namespaces and security configurations:

```

1 # kubeflow-namespaces.yaml
2 apiVersion: v1
3 kind: Namespace
4 metadata:
5   name: kubeflow
6   labels:
7     control-plane: kubeflow
8     istio-injection: enabled
9 ---
10 apiVersion: v1
11 kind: Namespace
12 metadata:
13   name: kubeflow-user-example-com
14   labels:
15     control-plane: kubeflow
16     istio-injection: enabled
17     user: example@example.com
18 ---
19 apiVersion: v1
20 kind: ServiceAccount
21 metadata:
22   name: kubeflow-admin
23   namespace: kubeflow
24 ---
25 apiVersion: rbac.authorization.k8s.io/v1
26 kind: ClusterRoleBinding
27 metadata:
28   name: kubeflow-admin
29 roleRef:

```

```

30   apiGroup: rbac.authorization.k8s.io
31   kind: ClusterRole
32   name: cluster-admin
33 subjects:
34 - kind: ServiceAccount
35   name: kubeflow-admin
36   namespace: kubeflow

```

Listing 2: Namespace and RBAC Configuration

2.2.2 Kubeflow Manifests Installation

Install Kubeflow using the official manifests with customizations for production environments:

```

1  #!/bin/bash
2
3  set -e
4
5  # Configuration variables
6  export KF_VERSION="v1.7.0"
7  export KF_NAME="mlops-platform"
8  export BASE_DIR="${HOME}/kubeflow"
9  export KF_DIR="${BASE_DIR}/${KF_NAME}"
10 export CONFIG_URI="https://raw.githubusercontent.com/kubeflow/manifests/${KF_VERSION}/kfdef/kfctl_k8s_istio.v1.7.0.yaml"
11
12 # Create directory structure
13 mkdir -p ${KF_DIR}
14 cd ${KF_DIR}
15
16 # Download Kubeflow manifests
17 echo "Downloading Kubeflow manifests..."
18 wget -O kubeflow-manifests.tar.gz \
19   "https://github.com/kubeflow/manifests/archive/${KF_VERSION}.tar.gz"
20 tar -xzf kubeflow-manifests.tar.gz
21 cd manifests-${KF_VERSION#v}
22
23 # Apply custom configurations
24 echo "Applying custom configurations..."
25 cat > custom-config.yaml << EOF
26 apiVersion: kfdef.apps.kubeflow.org/v1
27 kind: KfDef
28 metadata:
29   name: ${KF_NAME}
30   namespace: kubeflow
31 spec:
32   applications:
33   - kustomizeConfig:
34     repoRef:
35       name: manifests
36       path: stacks/kubernetes/application/istio-1-16
37     name: istio-1-16
38   - kustomizeConfig:
39     repoRef:
40       name: manifests
41       path: stacks/kubernetes/application/cluster-local-gateway-1-16
42     name: cluster-local-gateway-1-16
43   - kustomizeConfig:
44     repoRef:
45       name: manifests
46       path: apps/pipeline/upstream/env/cert-manager/platform-agnostic-multi-user
47     name: kubeflow-pipelines

```

```

48 - kustomizeConfig:
49   repoRef:
50     name: manifests
51     path: apps/jupyter/jupyter-web-app/upstream/overlays/istio
52   name: jupyter-web-app
53 - kustomizeConfig:
54   repoRef:
55     name: manifests
56     path: apps/katib/upstream/installs/katib-with-kubeflow
57   name: katib
58 - kustomizeConfig:
59   repoRef:
60     name: manifests
61     path: apps/training-operator/upstream/overlays/kubeflow
62   name: training-operator
63 repos:
64 - name: manifests
65   uri: https://github.com/kubeflow/manifests/archive/${KF_VERSION}.tar.gz
66 EOF
67
68 # Install cert-manager first (required for webhooks)
69 echo "Installing cert-manager..."
70 kubectl apply -f https://github.com/cert-manager/cert-manager/releases/download
71   /v1.12.0/cert-manager.yaml
72 kubectl wait --for=condition=ready pod -l app=cert-manager -n cert-manager --
73   timeout=300s
74
75 # Install KubeFlow components
76 echo "Installing KubeFlow components..."
77 while ! kustomize build example | kubectl apply -f -; do
78   echo "Retrying to apply resources"
79   sleep 10
80 done
81
82 # Wait for deployments to be ready
83 echo "Waiting for KubeFlow components to be ready..."
84 kubectl wait --for=condition=ready pod -l app=istiod -n istio-system --timeout
85   =300s
86 kubectl wait --for=condition=ready pod -l app=istio-proxy -n kubeflow --timeout
87   =300s
88
89 echo "KubeFlow installation completed successfully!"

```

Listing 3: KubeFlow Installation Process

2.2.3 KubeFlow Post-Installation Configuration

Configure KubeFlow for production use with proper security and resource management:

```

1 # kubeflow-production-config.yaml
2 apiVersion: v1
3 kind: ConfigMap
4 metadata:
5   name: kubeflow-config
6   namespace: kubeflow
7 data:
8   # Pipeline configuration
9   pipeline-config.yaml: |
10     executorImage: gcr.io/ml-pipeline/api-server:2.0.0
11     cacheEnabled: true
12     cacheDatabase:
13       host: mysql.kubeflow.svc.cluster.local
14       port: "3306"

```



```

15     database: mlpipeline
16   objectStore:
17     endpoint: minio-service.kubeflow.svc.cluster.local:9000
18     bucket: mlpipeline
19     accessKey: minio
20     secretKey: minio123
21
22   # Resource limits
23   resource-limits.yaml: |
24     defaultRequests:
25       cpu: "100m"
26       memory: "128Mi"
27     defaultLimits:
28       cpu: "1000m"
29       memory: "1Gi"
30     maxRequests:
31       cpu: "8000m"
32       memory: "16Gi"
33 ---
34 apiVersion: v1
35 kind: Secret
36 metadata:
37   name: kubeflow-secrets
38   namespace: kubeflow
39 type: Opaque
40 data:
41   # Base64 encoded values
42   mysql-password: bWxwaXB1bGluZQ== # mlpipeline
43   minio-access-key: bWluaW8= # minio
44   minio-secret-key: bWluaW8xMjM= # minio123

```

Listing 4: Production Kubeflow Configuration

2.3 MLflow Setup and Integration

MLflow serves as the experiment tracking and model registry backbone of our MLOps platform. This section covers both standalone MLflow deployment and integration with Kubeflow.

2.3.1 MLflow Server Deployment

Deploy MLflow server with persistent storage and database backend:

```

1 # mlflow-deployment.yaml
2 apiVersion: apps/v1
3 kind: Deployment
4 metadata:
5   name: mlflow-server
6   namespace: kubeflow
7   labels:
8     app: mlflow-server
9 spec:
10   replicas: 2
11   selector:
12     matchLabels:
13       app: mlflow-server
14   template:
15     metadata:
16       labels:
17         app: mlflow-server
18     spec:
19       containers:
20         - name: mlflow-server

```

```

21     image: python:3.9-slim
22     ports:
23     - containerPort: 5000
24     env:
25     - name: MLFLOW_S3_ENDPOINT_URL
26       value: "http://minio-service.kubeflow.svc.cluster.local:9000"
27     - name: AWS_ACCESS_KEY_ID
28       valueFrom:
29         secretKeyRef:
30           name: mlflow-secrets
31           key: aws-access-key-id
32     - name: AWS_SECRET_ACCESS_KEY
33       valueFrom:
34         secretKeyRef:
35           name: mlflow-secrets
36           key: aws-secret-access-key
37     - name: MLFLOW_DATABASE_URI
38       valueFrom:
39         secretKeyRef:
40           name: mlflow-secrets
41           key: database-uri
42     command:
43     - /bin/bash
44     - -c
45     - |
46       pip install mlflow[extras]==2.7.1 psycpg2-binary boto3
47       mlflow server \
48         --host 0.0.0.0 \
49         --port 5000 \
50         --backend-store-uri ${MLFLOW_DATABASE_URI} \
51         --default-artifact-root s3://mlflow-artifacts/ \
52         --serve-artifacts
53     resources:
54       requests:
55       cpu: 100m
56       memory: 256Mi
57       limits:
58       cpu: 500m
59       memory: 1Gi
60     livenessProbe:
61       httpGet:
62       path: /health
63       port: 5000
64       initialDelaySeconds: 30
65       periodSeconds: 10
66     readinessProbe:
67       httpGet:
68       path: /health
69       port: 5000
70       initialDelaySeconds: 5
71       periodSeconds: 5
72 ---
73 apiVersion: v1
74 kind: Service
75 metadata:
76   name: mlflow-server
77   namespace: kubeflow
78 spec:
79   selector:
80     app: mlflow-server
81   ports:
82   - port: 5000
83     targetPort: 5000

```

```

84     protocol: TCP
85     type: ClusterIP
86 ---
87 apiVersion: v1
88 kind: Secret
89 metadata:
90   name: mlflow-secrets
91   namespace: kubeflow
92 type: Opaque
93 data:
94   aws-access-key-id: bWluaW8= # minio
95   aws-secret-access-key: bWluaW8xMjM= # minio123
96   database-uri:
     cG9zdGdyZXNxbDovL21sZmxvdzptbGZsb3dAcG9zdGdyZXNxbC5rdWJlZmxvdy5zdmMuY2x1c3Rlci5sb2NhbDo1N
     ==

```

Listing 5: MLflow Server Deployment

2.3.2 Database and Storage Configuration

Set up PostgreSQL database and MinIO object storage for MLflow:

```

1 # postgresql-deployment.yaml
2 apiVersion: apps/v1
3 kind: Deployment
4 metadata:
5   name: postgresql
6   namespace: kubeflow
7 spec:
8   replicas: 1
9   selector:
10     matchLabels:
11       app: postgresql
12   template:
13     metadata:
14       labels:
15         app: postgresql
16     spec:
17       containers:
18       - name: postgresql
19         image: postgres:13
20         env:
21         - name: POSTGRES_DB
22           value: mlflow
23         - name: POSTGRES_USER
24           value: mlflow
25         - name: POSTGRES_PASSWORD
26           value: mlflow
27         - name: PGDATA
28           value: /var/lib/postgresql/data/pgdata
29       ports:
30       - containerPort: 5432
31       volumeMounts:
32       - name: postgresql-storage
33         mountPath: /var/lib/postgresql/data
34       resources:
35         requests:
36           cpu: 100m
37           memory: 256Mi
38         limits:
39           cpu: 500m
40           memory: 1Gi
41       volumes:

```

```

42     - name: postgresql-storage
43       persistentVolumeClaim:
44         claimName: postgresql-pvc
45 ---
46 apiVersion: v1
47 kind: Service
48 metadata:
49   name: postgresql
50   namespace: kubeflow
51 spec:
52   selector:
53     app: postgresql
54   ports:
55     - port: 5432
56       targetPort: 5432
57 ---
58 apiVersion: v1
59 kind: PersistentVolumeClaim
60 metadata:
61   name: postgresql-pvc
62   namespace: kubeflow
63 spec:
64   accessModes:
65     - ReadWriteOnce
66   resources:
67     requests:
68       storage: 10Gi
69   storageClassName: fast-ssd
70 ---
71 # minio-deployment.yaml
72 apiVersion: apps/v1
73 kind: Deployment
74 metadata:
75   name: minio
76   namespace: kubeflow
77 spec:
78   replicas: 1
79   selector:
80     matchLabels:
81       app: minio
82   template:
83     metadata:
84       labels:
85         app: minio
86     spec:
87       containers:
88         - name: minio
89           image: minio/minio:RELEASE.2023-09-04T19-57-37Z
90           args:
91             - server
92             - /data
93             - --console-address=:9001
94           env:
95             - name: MINIO_ROOT_USER
96               value: minio
97             - name: MINIO_ROOT_PASSWORD
98               value: minio123
99           ports:
100             - containerPort: 9000
101             - containerPort: 9001
102           volumeMounts:
103             - name: minio-storage
104               mountPath: /data

```

```

105     resources:
106         requests:
107             cpu: 100m
108             memory: 256Mi
109         limits:
110             cpu: 500m
111             memory: 1Gi
112     volumes:
113     - name: minio-storage
114       persistentVolumeClaim:
115         claimName: minio-pvc
116 ---
117 apiVersion: v1
118 kind: Service
119 metadata:
120     name: minio-service
121     namespace: kubeflow
122 spec:
123     selector:
124         app: minio
125     ports:
126     - name: api
127       port: 9000
128       targetPort: 9000
129     - name: console
130       port: 9001
131       targetPort: 9001
132 ---
133 apiVersion: v1
134 kind: PersistentVolumeClaim
135 metadata:
136     name: minio-pvc
137     namespace: kubeflow
138 spec:
139     accessModes:
140     - ReadWriteOnce
141     resources:
142         requests:
143             storage: 50Gi
144     storageClassName: fast-ssd

```

Listing 6: MLflow Storage Infrastructure

2.4 Network Configuration and Security

Proper network configuration ensures secure communication between components while maintaining performance and accessibility.

2.4.1 Istio Service Mesh Configuration

Configure Istio for secure service-to-service communication:

```

1 # istio-security-policies.yaml
2 apiVersion: security.istio.io/v1beta1
3 kind: PeerAuthentication
4 metadata:
5     name: default
6     namespace: kubeflow
7 spec:
8     mtls:
9         mode: STRICT
10 ---

```

```

11 apiVersion: security.istio.io/v1beta1
12 kind: AuthorizationPolicy
13 metadata:
14   name: mlflow-access
15   namespace: kubeflow
16 spec:
17   selector:
18     matchLabels:
19       app: mlflow-server
20   rules:
21   - from:
22     - source:
23       principals: ["cluster.local/ns/kubeflow/sa/default"]
24     - source:
25       namespaces: ["kubeflow-user-example-com"]
26   - to:
27     - operation:
28       methods: ["GET", "POST", "PUT", "DELETE"]
29 ---
30 apiVersion: networking.istio.io/v1beta1
31 kind: VirtualService
32 metadata:
33   name: mlflow-vs
34   namespace: kubeflow
35 spec:
36   hosts:
37   - mlflow.example.com
38   gateways:
39   - kubeflow-gateway
40   http:
41   - match:
42     - uri:
43       prefix: /
44     route:
45     - destination:
46       host: mlflow-server.kubeflow.svc.cluster.local
47       port:
48         number: 5000
49     timeout: 300s
50 ---
51 apiVersion: networking.istio.io/v1beta1
52 kind: Gateway
53 metadata:
54   name: kubeflow-gateway
55   namespace: kubeflow
56 spec:
57   selector:
58     istio: ingressgateway
59   servers:
60   - port:
61     number: 80
62     name: http
63     protocol: HTTP
64     hosts:
65     - "*.example.com"
66     tls:
67       httpsRedirect: true
68   - port:
69     number: 443
70     name: https
71     protocol: HTTPS
72     hosts:
73     - "*.example.com"

```

```

74     tls:
75       mode: SIMPLE
76       credentialName: kubeflow-tls-secret

```

Listing 7: Istio Security Configuration

2.5 Monitoring and Observability Setup

Comprehensive monitoring is essential for maintaining platform health and performance.

2.5.1 Prometheus and Grafana Configuration

Deploy monitoring stack with custom dashboards for MLOps metrics:

```

1 # monitoring-namespace.yaml
2 apiVersion: v1
3 kind: Namespace
4 metadata:
5   name: monitoring
6   labels:
7     name: monitoring
8 ---
9 # prometheus-config.yaml
10 apiVersion: v1
11 kind: ConfigMap
12 metadata:
13   name: prometheus-config
14   namespace: monitoring
15 data:
16   prometheus.yml: |
17     global:
18       scrape_interval: 15s
19       evaluation_interval: 15s
20
21     rule_files:
22       - "mlops_rules.yml"
23
24     scrape_configs:
25       - job_name: 'kubeflow-pipelines'
26         kubernetes_sd_configs:
27           - role: pod
28             namespaces:
29               names:
30                 - kubeflow
31         relabel_configs:
32           - source_labels: [__meta_kubernetes_pod_label_app]
33             action: keep
34             regex: ml-pipeline.*
35
36       - job_name: 'mlflow-server'
37         kubernetes_sd_configs:
38           - role: pod
39             namespaces:
40               names:
41                 - kubeflow
42         relabel_configs:
43           - source_labels: [__meta_kubernetes_pod_label_app]
44             action: keep
45             regex: mlflow-server
46
47       - job_name: 'model-servers'
48         kubernetes_sd_configs:
49           - role: pod

```

```

50     namespaces:
51       names:
52         - kubeflow-user-example-com
53     relabel_configs:
54       - source_labels: [
55         __meta_kubernetes_pod_annotation_serving_kserve_io_inferenceservice]
56         action: keep
57         regex: .+
58 mlops_rules.yml: |
59   groups:
60     - name: mlops.rules
61       rules:
62         - alert: ModelServerDown
63           expr: up{job="model-servers"} == 0
64           for: 1m
65           labels:
66             severity: critical
67           annotations:
68             summary: "Model server {{ $labels.instance }} is down"
69             description: "Model server has been down for more than 1 minute"
70
71         - alert: HighModelLatency
72           expr: histogram_quantile(0.95, rate(
73 model_request_duration_seconds_bucket[5m])) > 1
74           for: 2m
75           labels:
76             severity: warning
77           annotations:
78             summary: "High model inference latency detected"
79             description: "95th percentile latency is {{ $value }}s"
80
81         - alert: MLflowServerDown
82           expr: up{job="mlflow-server"} == 0
83           for: 2m
84           labels:
85             severity: critical
86           annotations:
87             summary: "MLflow server is down"
88             description: "MLflow server has been unreachable for more than 2
minutes"

```

Listing 8: Monitoring Stack Deployment

2.6 Initial Platform Validation

After completing the installation, validate that all components are functioning correctly:

```

1  #!/bin/bash
2
3  set -e
4
5  echo "=== MLOps Platform Validation ==="
6
7  # Check namespace status
8  echo "Checking namespace status..."
9  kubectl get namespaces kubeflow kubeflow-user-example-com monitoring
10
11 # Check Kubeflow components
12 echo "Checking Kubeflow components..."
13 kubectl get pods -n kubeflow | grep -E "(Running|Completed)" || exit 1
14
15 # Check MLflow server

```



```

16 echo "Checking MLflow server..."
17 kubectl get pods -n kubeflow -l app=mlflow-server
18 kubectl wait --for=condition=ready pod -l app=mlflow-server -n kubeflow --
    timeout=300s
19
20 # Test MLflow API
21 echo "Testing MLflow API..."
22 MLF_HOST=$(kubectl get svc mlflow-server -n kubeflow -o jsonpath='{.spec.
    clusterIP}')
23 kubectl run test-pod --rm -i --tty --image=curlimages/curl -- \
24     curl -f http://${MLF_HOST}:5000/health || exit 1
25
26 # Check storage components
27 echo "Checking storage components..."
28 kubectl get pods -n kubeflow -l app=postgresql
29 kubectl get pods -n kubeflow -l app=minio
30
31 # Check Istio configuration
32 echo "Checking Istio configuration..."
33 kubectl get virtualservices,gateways -n kubeflow
34
35 # Test pipeline functionality
36 echo "Testing pipeline functionality..."
37 kubectl get workflows -n kubeflow-user-example-com || echo "No workflows found
    (expected for new installation)"
38
39 echo "=== Platform validation completed successfully! ==="
40 echo ""
41 echo "Access URLs (configure DNS or port-forward):"
42 echo "- Kubeflow Central Dashboard: https://kubeflow.example.com"
43 echo "- MLflow UI: https://mlflow.example.com"
44 echo "- MinIO Console: https://minio.example.com"
45 echo ""
46 echo "Next steps:"
47 echo "1. Configure DNS entries for your domain"
48 echo "2. Set up SSL certificates"
49 echo "3. Create user profiles and RBAC policies"
50 echo "4. Run your first ML pipeline"

```

Listing 9: Platform Validation Script

This implementation guide provides a solid foundation for deploying a production-ready MLOps platform. The next sections will cover advanced pipeline development, model serving configurations, and operational best practices.

3 Model Deployment and Serving

Model deployment and serving represent critical phases in the MLOps lifecycle where trained models transition from experimental artifacts to production systems that deliver real business value. This section provides comprehensive guidance for implementing robust, scalable, and secure model serving infrastructure using KServe, along with advanced deployment patterns that ensure safe model releases and optimal performance.

3.1 KServe Architecture and Integration

KServe (formerly KFServing) serves as the cornerstone of our model serving strategy, providing a Kubernetes-native platform for deploying and managing machine learning models at scale. Built on the foundation of Knative Serving, KServe inherits powerful capabilities for autoscaling, traffic management, and serverless deployment patterns while adding ML-specific functionality.

3.1.1 KServe Core Components

Understanding KServe's architecture is essential for effective model deployment and troubleshooting:

InferenceService: The primary custom resource that defines how models should be served, including predictor configurations, transformers, and explainers.

Model Server: Framework-specific serving runtimes that handle model loading, inference request processing, and response formatting. KServe supports multiple built-in servers including:

- **Scikit-learn Server:** Optimized for traditional ML models with pickle format support
- **TensorFlow Serving:** High-performance serving for TensorFlow models with batching and GPU support
- **PyTorch Server (TorchServe):** Native PyTorch model serving with custom handler support
- **XGBoost Server:** Specialized serving for gradient boosting models
- **Custom Predictors:** User-defined serving logic for complex inference pipelines

Data Plane: Handles actual inference requests and responses, implementing the KServe v1 and v2 inference protocols for standardized communication.

Control Plane: Manages the lifecycle of inference services, including deployment, scaling, and traffic routing decisions.

3.1.2 KServe Installation and Configuration

Deploy KServe with production-ready configurations:

```

1  #!/bin/bash
2
3  set -e
4
5  echo "Installing KServe with dependencies..."
6
7  # Install Knative Serving (required for KServe)
8  echo "Installing Knative Serving..."
9  kubectl apply -f https://github.com/knative/serving/releases/download/knative-
    v1.11.0/serving-crds.yaml
10 kubectl apply -f https://github.com/knative/serving/releases/download/knative-
    v1.11.0/serving-core.yaml
11
12 # Install Knative Istio controller
13 kubectl apply -f https://github.com/knative/net-istio/releases/download/knative
    -v1.11.0/net-istio.yaml
14
15 # Configure Knative Serving
16 kubectl patch configmap/config-network \
17   --namespace knative-serving \
18   --type merge \
19   --patch '{"data":{"ingress-class":"istio.ingress.networking.knative.dev"}}'
20
21 # Set domain configuration
22 kubectl patch configmap/config-domain \
23   --namespace knative-serving \
24   --type merge \
25   --patch '{"data":{"example.com":""}}'
26
27 # Install KServe CRDs and controllers

```

```

28 echo "Installing KServe..."
29 kubectl apply -f https://github.com/kserve/kserve/releases/download/v0.11.0/
    kserve.yaml
30
31 # Install KServe built-in ClusterServingRuntimes
32 kubectl apply -f https://github.com/kserve/kserve/releases/download/v0.11.0/
    kserve-runtimes.yaml
33
34 # Wait for KServe controller to be ready
35 echo "Waiting for KServe controller..."
36 kubectl wait --for=condition=ready pod -l control-plane=kserve-controller-
    manager -n kserve --timeout=300s
37
38 echo "KServe installation completed successfully!"

```

Listing 10: KServe Installation Script

3.1.3 Production KServe Configuration

Configure KServe for production environments with proper resource management and security:

```

1 # kserve-config.yaml
2 apiVersion: v1
3 kind: ConfigMap
4 metadata:
5   name: inferencesservice-config
6   namespace: kserve
7 data:
8   predictors: |
9     {
10       "tensorflow": {
11         "image": "tensorflow/serving:2.13.0",
12         "defaultImageVersion": "2.13.0",
13         "defaultGpuImageVersion": "2.13.0-gpu",
14         "supportedFrameworks": ["tensorflow"],
15         "multiModelServer": false
16       },
17       "pytorch": {
18         "image": "pytorch/torchserve:0.8.2-cpu",
19         "defaultImageVersion": "0.8.2-cpu",
20         "defaultGpuImageVersion": "0.8.2-gpu",
21         "supportedFrameworks": ["pytorch"],
22         "multiModelServer": false
23       },
24       "sklearn": {
25         "image": "kserve/sklearnserver:v0.11.0",
26         "defaultImageVersion": "v0.11.0",
27         "supportedFrameworks": ["sklearn"],
28         "multiModelServer": true
29       },
30       "xgboost": {
31         "image": "kserve/xgbserver:v0.11.0",
32         "defaultImageVersion": "v0.11.0",
33         "supportedFrameworks": ["xgboost"],
34         "multiModelServer": true
35       }
36     }
37
38   transformer: |
39     {
40       "feast": {
41         "image": "kserve/feast-transformer:v0.11.0",
42         "defaultImageVersion": "v0.11.0"

```

```

43     }
44   }
45
46   explainer: |
47     {
48       "alibi": {
49         "image": "kserve/alibi-explainer:v0.11.0",
50         "defaultImageVersion": "v0.11.0"
51       }
52     }
53
54   storageInitializer: |
55     {
56       "image": "kserve/storage-initializer:v0.11.0",
57       "memoryRequest": "100Mi",
58       "memoryLimit": "1Gi",
59       "cpuRequest": "100m",
60       "cpuLimit": "1000m"
61     }
62
63   credentials: |
64     {
65       "gcs": {
66         "gcsCredentialFileName": "gcloud-application-credentials.json"
67       },
68       "s3": {
69         "s3AccessKeyIdName": "AWS_ACCESS_KEY_ID",
70         "s3SecretAccessKeyName": "AWS_SECRET_ACCESS_KEY",
71         "s3Endpoint": "",
72         "s3UseHttps": true,
73         "s3Region": "us-west-1",
74         "s3VerifySSL": true,
75         "s3UseVirtualBucket": false,
76         "s3UseAnonymousCredential": false,
77         "s3CABundle": ""
78       }
79     }
80
81   ingress: |
82     {
83       "ingressGateway": "kubeflow/kubeflow-gateway",
84       "ingressService": "istio-ingressgateway.istio-system.svc.cluster.local",
85       "localGateway": "knative-serving/knative-local-gateway",
86       "localGatewayService": "knative-local-gateway.istio-system.svc.cluster.
87       local",
88       "ingressDomain": "example.com",
89       "ingressClassName": "istio",
90       "domainTemplate": "{{.Name}}-{{.Namespace}}.{{.IngressDomain}}",
91       "urlScheme": "https",
92       "disableIstioVirtualHost": false
93     }
94
95   deploy: |
96     {
97       "defaultDeploymentMode": "Serverless",
98       "progressDeadlineSeconds": 600,
99       "defaultCpuRequest": "100m",
100      "defaultMemoryRequest": "128Mi",
101      "defaultCpuLimit": "1000m",
102      "defaultMemoryLimit": "2Gi"
103    }
104 ---
105 apiVersion: v1

```

```

105 kind: ConfigMap
106 metadata:
107   name: kserve-logger-config
108   namespace: kserve
109 data:
110   logger.properties: |
111     # Root logger option
112     log4j.rootLogger=INFO, stdout
113
114     # Direct log messages to stdout
115     log4j.appender.stdout=org.apache.log4j.ConsoleAppender
116     log4j.appender.stdout.Target=System.out
117     log4j.appender.stdout.layout=org.apache.log4j.PatternLayout
118     log4j.appender.stdout.layout.ConversionPattern=%d{yyyy-MM-dd HH:mm:ss} %-5p
119       %c{1}:%L - %m%n
120
121     # Suppress unnecessary logs
122     log4j.logger.org.apache.hadoop=WARN
123     log4j.logger.org.apache.spark=WARN
124     log4j.logger.org.eclipse.jetty=WARN
125     log4j.logger.org.apache.kafka=WARN

```

Listing 11: KServe Production Configuration

3.2 Model Packaging and Registry Integration

Effective model deployment requires standardized packaging and integration with model registries to ensure consistency, traceability, and reproducibility.

3.2.1 MLflow Model Integration

Create a comprehensive MLflow integration that automatically packages models for KServe deployment:

```

1 import mlflow
2 import mlflow.sklearn
3 import mlflow.pytorch
4 import mlflow.tensorflow
5 import yaml
6 import json
7 import os
8 from typing import Dict, Any, Optional
9 from kubernetes import client, config
10 from datetime import datetime
11 import logging
12
13 class MLflowKServeIntegration:
14     """
15     Handles integration between MLflow model registry and KServe deployments
16     """
17
18     def __init__(self,
19                 mlflow_uri: str,
20                 namespace: str = "kubeflow-user-example-com",
21                 kserve_domain: str = "example.com"):
22         self.mlflow_uri = mlflow_uri
23         self.namespace = namespace
24         self.kserve_domain = kserve_domain
25
26         # Initialize MLflow client
27         mlflow.set_tracking_uri(mlflow_uri)
28         self.mlflow_client = mlflow.MlflowClient()

```

```

29
30     # Initialize Kubernetes client
31     try:
32         config.load_incluster_config()
33     except:
34         config.load_kube_config()
35
36     self.k8s_client = client.ApiClient()
37     self.custom_client = client.CustomObjectsApi()
38
39     # Configure logging
40     logging.basicConfig(level=logging.INFO)
41     self.logger = logging.getLogger(__name__)
42
43     def get_model_info(self, model_name: str, version: Optional[str] = None) ->
44     Dict[str, Any]:
45         """
46         Retrieve comprehensive model information from MLflow registry
47         """
48         try:
49             if version is None:
50                 # Get latest version
51                 latest_versions = self.mlflow_client.get_latest_versions(
52                     model_name, stages=["Production", "Staging"]
53                 )
54                 if not latest_versions:
55                     raise ValueError(f"No versions found for model {model_name}")
56
57                 model_version = latest_versions[0]
58             else:
59                 model_version = self.mlflow_client.get_model_version(model_name
60 , version)
61
62             # Get run information
63             run = self.mlflow_client.get_run(model_version.run_id)
64
65             # Extract model artifacts
66             model_uri = f"models://{model_name}/{model_version.version}"
67
68             return {
69                 "name": model_name,
70                 "version": model_version.version,
71                 "stage": model_version.current_stage,
72                 "model_uri": model_uri,
73                 "run_id": model_version.run_id,
74                 "framework": self._detect_framework(run),
75                 "artifacts": run.data.tags.get("mlflow.log-model.history", "[]")
76             },
77
78             "metrics": run.data.metrics,
79             "params": run.data.params,
80             "tags": run.data.tags,
81             "creation_timestamp": model_version.creation_timestamp,
82             "last_updated_timestamp": model_version.last_updated_timestamp
83         }
84
85     except Exception as e:
86         self.logger.error(f"Error retrieving model info: {str(e)}")
87         raise
88
89     def _detect_framework(self, run) -> str:
90         """
91         Detect ML framework from run artifacts and tags
92         """

```

```

88     artifacts = json.loads(run.data.tags.get("mlflow.log-model.history", "
    []"))
89
90     for artifact in artifacts:
91         if "sklearn" in artifact.get("flavors", {}):
92             return "sklearn"
93         elif "pytorch" in artifact.get("flavors", {}):
94             return "pytorch"
95         elif "tensorflow" in artifact.get("flavors", {}):
96             return "tensorflow"
97         elif "xgboost" in artifact.get("flavors", {}):
98             return "xgboost"
99
100     # Fallback to run tags
101     if "mlflow.source.type" in run.data.tags:
102         source_type = run.data.tags["mlflow.source.type"].lower()
103         if "sklearn" in source_type:
104             return "sklearn"
105         elif "pytorch" in source_type:
106             return "pytorch"
107         elif "tensorflow" in source_type:
108             return "tensorflow"
109
110     return "sklearn" # Default fallback
111
112     def create_inference_service(self,
113                                model_name: str,
114                                model_version: Optional[str] = None,
115                                service_name: Optional[str] = None,
116                                resources: Optional[Dict[str, Any]] = None,
117                                autoscaling: Optional[Dict[str, Any]] = None,
118                                canary_percent: int = 0) -> Dict[str, Any]:
119         """
120         Create KServe InferenceService from MLflow model
121         """
122
123         # Get model information
124         model_info = self.get_model_info(model_name, model_version)
125
126         if service_name is None:
127             service_name = f"{model_name.lower().replace('_', '-')}-v{
model_info['version']}"
128
129         # Default resource configuration
130         default_resources = {
131             "requests": {"cpu": "100m", "memory": "256Mi"},
132             "limits": {"cpu": "1000m", "memory": "2Gi"}
133         }
134         if resources:
135             default_resources.update(resources)
136
137         # Default autoscaling configuration
138         default_autoscaling = {
139             "minReplicas": 1,
140             "maxReplicas": 10,
141             "targetUtilizationPercentage": 70,
142             "scaleToZeroGracePeriod": "30s",
143             "scaleDownDelay": "0s",
144             "stableWindow": "60s"
145         }
146         if autoscaling:
147             default_autoscaling.update(autoscaling)
148

```

```

149     # Create InferenceService specification
150     inference_service = {
151         "apiVersion": "serving.kserve.io/v1beta1",
152         "kind": "InferenceService",
153         "metadata": {
154             "name": service_name,
155             "namespace": self.namespace,
156             "labels": {
157                 "model-name": model_name,
158                 "model-version": str(model_info['version']),
159                 "model-stage": model_info['stage'],
160                 "framework": model_info['framework'],
161                 "managed-by": "mlflow-kserve-integration"
162             },
163             "annotations": {
164                 "mlflow.model.uri": model_info['model_uri'],
165                 "mlflow.run.id": model_info['run_id'],
166                 "deployment.timestamp": datetime.utcnow().isoformat(),
167                 "serving.kserve.io/deploymentMode": "Serverless"
168             }
169         },
170         "spec": {
171             "predictor": {
172                 model_info['framework']: {
173                     "storageUri": model_info['model_uri'],
174                     "resources": default_resources,
175                     "env": [
176                         {
177                             "name": "STORAGE_URI",
178                             "value": model_info['model_uri']
179                         },
180                         {
181                             "name": "MODEL_NAME",
182                             "value": model_name
183                         }
184                     ]
185                 }
186             }
187         }
188     }
189
190     # Add canary configuration if specified
191     if canary_percent > 0:
192         inference_service["spec"]["predictor"]["canaryTrafficPercent"] =
canary_percent
193
194     # Add autoscaling annotations
195     inference_service["metadata"]["annotations"].update({
196         "autoscaling.knative.dev/minScale": str(default_autoscaling["
minReplicas"]),
197         "autoscaling.knative.dev/maxScale": str(default_autoscaling["
maxReplicas"]),
198         "autoscaling.knative.dev/target": str(default_autoscaling["
targetUtilizationPercentage"]),
199         "autoscaling.knative.dev/scaleToZeroGracePeriod":
default_autoscaling["scaleToZeroGracePeriod"],
200         "autoscaling.knative.dev/scaleDownDelay": default_autoscaling["
scaleDownDelay"],
201         "autoscaling.knative.dev/window": default_autoscaling["stableWindow
"]
202     })
203
204     # Deploy to Kubernetes

```



```

205         try:
206             response = self.custom_client.create_namespaced_custom_object(
207                 group="serving.kserve.io",
208                 version="v1beta1",
209                 namespace=self.namespace,
210                 plural="inferenceservices",
211                 body=inference_service
212             )
213
214             self.logger.info(f"Successfully created InferenceService: {
service_name}")
215
216             # Update MLflow model stage if deploying to production
217             if model_info['stage'] != "Production":
218                 self.mlflow_client.transition_model_version_stage(
219                     name=model_name,
220                     version=model_info['version'],
221                     stage="Production",
222                     archive_existing_versions=False
223                 )
224
225             return {
226                 "service_name": service_name,
227                 "namespace": self.namespace,
228                 "model_info": model_info,
229                 "inference_service": response,
230                 "endpoint_url": f"https://{service_name}-{self.namespace}.{self
.kserve_domain}"
231             }
232
233         except Exception as e:
234             self.logger.error(f"Failed to create InferenceService: {str(e)}")
235             raise
236
237     def update_traffic_split(self,
238                             service_name: str,
239                             traffic_config: Dict[str, int]) -> Dict[str, Any]:
240
241         """
242         Update traffic splitting between model versions
243         """
244         try:
245             # Get current InferenceService
246             current_service = self.custom_client.get_namespaced_custom_object(
247                 group="serving.kserve.io",
248                 version="v1beta1",
249                 namespace=self.namespace,
250                 plural="inferenceservices",
251                 name=service_name
252             )
253
254             # Update traffic configuration
255             if "canary" in traffic_config:
256                 patch_body = {
257                     "spec": {
258                         "predictor": {
259                             "canaryTrafficPercent": traffic_config["canary"]
260                         }
261                     }
262                 }
263
264             response = self.custom_client.patch_namespaced_custom_object(
265                 group="serving.kserve.io",
266                 version="v1beta1",

```

```

266         namespace=self.namespace,
267         plural="inferenceservices",
268         name=service_name,
269         body=patch_body
270     )
271
272     self.logger.info(f"Updated traffic split for {service_name}: {
traffic_config}")
273     return response
274
275     except Exception as e:
276         self.logger.error(f"Failed to update traffic split: {str(e)}")
277         raise
278
279     def rollback_deployment(self, service_name: str, target_version: str) ->
Dict[str, Any]:
280         """
281         Rollback deployment to a previous model version
282         """
283         try:
284             # Get service metadata to extract model info
285             current_service = self.custom_client.get_namespaced_custom_object(
286                 group="serving.kserve.io",
287                 version="v1beta1",
288                 namespace=self.namespace,
289                 plural="inferenceservices",
290                 name=service_name
291             )
292
293             model_name = current_service["metadata"]["labels"]["model-name"]
294
295             # Get target version info
296             target_model_info = self.get_model_info(model_name, target_version)
297
298             # Update the service to use target version
299             patch_body = {
300                 "spec": {
301                     "predictor": {
302                         target_model_info['framework']: {
303                             "storageUri": target_model_info['model_uri']
304                         }
305                     }
306                 },
307                 "metadata": {
308                     "labels": {
309                         "model-version": str(target_model_info['version'])
310                     },
311                     "annotations": {
312                         "mlflow.model.uri": target_model_info['model_uri'],
313                         "mlflow.run.id": target_model_info['run_id'],
314                         "rollback.timestamp": datetime.utcnow().isoformat(),
315                         "rollback.target.version": target_version
316                     }
317                 }
318             }
319
320             response = self.custom_client.patch_namespaced_custom_object(
321                 group="serving.kserve.io",
322                 version="v1beta1",
323                 namespace=self.namespace,
324                 plural="inferenceservices",
325                 name=service_name,
326                 body=patch_body

```

```

327         )
328
329         self.logger.info(f"Successfully rolled back {service_name} to
version {target_version}")
330         return response
331
332     except Exception as e:
333         self.logger.error(f"Failed to rollback deployment: {str(e)}")
334         raise
335
336 # Usage example
337 if __name__ == "__main__":
338     # Initialize integration
339     integration = MLflowKServeIntegration(
340         mlflow_uri="http://mlflow-server.kubeflow.svc.cluster.local:5000",
341         namespace="kubeflow-user-example-com"
342     )
343
344     # Deploy a model
345     deployment_result = integration.create_inference_service(
346         model_name="fraud-detection-model",
347         resources={
348             "requests": {"cpu": "200m", "memory": "512Mi"},
349             "limits": {"cpu": "2000m", "memory": "4Gi"}
350         },
351         autoscaling={
352             "minReplicas": 2,
353             "maxReplicas": 20,
354             "targetUtilizationPercentage": 80
355         }
356     )
357
358     print(f"Model deployed successfully: {deployment_result['endpoint_url']}")

```

Listing 12: MLflow KServe Integration

3.3 Advanced Deployment Patterns

Production model deployment requires sophisticated strategies to minimize risk while ensuring continuous service availability. This section covers advanced deployment patterns that enable safe model releases.

3.3.1 Canary Deployments

Implement gradual rollout of new model versions with automated monitoring and rollback capabilities:

```

1 import asyncio
2 import logging
3 from typing import Dict, List, Optional
4 from dataclasses import dataclass
5 from datetime import datetime, timedelta
6 import numpy as np
7 from prometheus_client.parser import text_string_to_metric_families
8 import aiohttp
9 import yaml
10
11 @dataclass
12 class CanaryConfig:
13     """Configuration for canary deployment"""
14     initial_traffic_percent: int = 5
15     increment_percent: int = 10

```

```

16     max_traffic_percent: int = 50
17     evaluation_duration_minutes: int = 10
18     success_rate_threshold: float = 0.99
19     latency_threshold_ms: float = 1000
20     error_rate_threshold: float = 0.01
21     auto_promote: bool = True
22     auto_rollback: bool = True
23
24 @dataclass
25 class DeploymentMetrics:
26     """Metrics for deployment evaluation"""
27     success_rate: float
28     avg_latency_ms: float
29     error_rate: float
30     request_count: int
31     timestamp: datetime
32
33 class CanaryDeploymentController:
34     """
35     Automated canary deployment controller with monitoring and decision making
36     """
37
38     def __init__(self,
39                 kserve_integration: MLflowKServeIntegration,
40                 prometheus_url: str,
41                 config: CanaryConfig):
42         self.kserve_integration = kserve_integration
43         self.prometheus_url = prometheus_url
44         self.config = config
45         self.logger = logging.getLogger(__name__)
46
47     async def deploy_canary(self,
48                           model_name: str,
49                           new_version: str,
50                           service_name: str) -> Dict[str, Any]:
51         """
52         Execute automated canary deployment process
53         """
54
55         deployment_log = {
56             "model_name": model_name,
57             "new_version": new_version,
58             "service_name": service_name,
59             "start_time": datetime.utcnow(),
60             "stages": [],
61             "final_status": "in_progress"
62         }
63
64         try:
65             # Stage 1: Deploy new version with minimal traffic
66             self.logger.info(f"Starting canary deployment for {model_name} v{
new_version}")
67
68             # Create new InferenceService for canary version
69             canary_service_name = f"{service_name}-canary"
70             canary_deployment = self.kserve_integration.
create_inference_service(
71                 model_name=model_name,
72                 model_version=new_version,
73                 service_name=canary_service_name,
74                 canary_percent=self.config.initial_traffic_percent
75             )
76

```

```

77         deployment_log["stages"].append({
78             "stage": "initial_deployment",
79             "traffic_percent": self.config.initial_traffic_percent,
80             "timestamp": datetime.utcnow(),
81             "status": "success"
82         })
83
84         # Wait for deployment to stabilize
85         await asyncio.sleep(60)
86
87         # Stage 2: Gradual traffic increase with monitoring
88         current_traffic = self.config.initial_traffic_percent
89
90         while current_traffic < self.config.max_traffic_percent:
91             # Evaluate current performance
92             baseline_metrics = await self.get_deployment_metrics(
service_name)
93             canary_metrics = await self.get_deployment_metrics(
canary_service_name)
94
95             # Make deployment decision
96             decision = self._evaluate_canary_performance(baseline_metrics,
canary_metrics)
97
98             stage_log = {
99                 "stage": "traffic_increase",
100                 "traffic_percent": current_traffic,
101                 "timestamp": datetime.utcnow(),
102                 "baseline_metrics": baseline_metrics.__dict__ if
baseline_metrics
103                 else None,
104                 "canary_metrics": canary_metrics.__dict__ if canary_metrics
105                 else None,
106                 "decision": decision
107             }
108
109             if decision["action"] == "continue":
110                 # Increase traffic
111                 current_traffic = min(
112                     current_traffic + self.config.increment_percent,
113                     self.config.max_traffic_percent
114                 )
115
116                 await self._update_traffic_split(canary_service_name,
current_traffic)
117
118                 stage_log["new_traffic_percent"] = current_traffic
119                 stage_log["status"] = "success"
120
121                 self.logger.info(f"Increased canary traffic to {
current_traffic}%")
122
123                 elif decision["action"] == "rollback":
124                     # Automatic rollback
125                     if self.config.auto_rollback:
126                         await self._rollback_canary(service_name,
canary_service_name)
127
128                         stage_log["status"] = "rollback"
129                         deployment_log["final_status"] = "failed"
130                         break
131                     else:
132                         # Manual intervention required
133                         stage_log["status"] = "requires_manual_intervention"
134                         deployment_log["final_status"] = "requires_intervention"
135
136         "

```

```

131         break
132
133         elif decision["action"] == "hold":
134             # Hold current traffic level for extended evaluation
135             self.logger.info(f"Holding canary traffic at {
current_traffic}% for extended evaluation")
136             await asyncio.sleep(self.config.evaluation_duration_minutes
* 60 * 2) # Extended wait
137             stage_log["status"] = "hold"
138
139             deployment_log["stages"].append(stage_log)
140
141             # Wait for evaluation period
142             await asyncio.sleep(self.config.evaluation_duration_minutes *
60)
143
144             # Stage 3: Final evaluation and promotion decision
145             if deployment_log["final_status"] == "in_progress":
146                 final_baseline_metrics = await self.get_deployment_metrics(
service_name)
147                 final_canary_metrics = await self.get_deployment_metrics(
canary_service_name)
148
149                 final_decision = self._evaluate_canary_performance(
150                     final_baseline_metrics,
151                     final_canary_metrics,
152                     final_evaluation=True
153                 )
154
155                 if final_decision["action"] == "promote" and self.config.
auto_promote:
156                     # Promote canary to production
157                     await self._promote_canary(service_name,
canary_service_name, new_version)
158                     deployment_log["final_status"] = "promoted"
159                     self.logger.info(f"Successfully promoted {model_name} v{
new_version} to production")
160                 else:
161                     deployment_log["final_status"] = "requires_manual_promotion
"
162                     self.logger.info(f"Canary deployment ready for manual
promotion")
163
164                     deployment_log["stages"].append({
165                         "stage": "final_evaluation",
166                         "timestamp": datetime.utcnow(),
167                         "final_decision": final_decision,
168                         "status": deployment_log["final_status"]
169                     })
170
171                     deployment_log["end_time"] = datetime.utcnow()
172                     deployment_log["duration_minutes"] = (
173                         deployment_log["end_time"] - deployment_log["start_time"]
174                     ).total_seconds() / 60
175
176                     return deployment_log
177
178             except Exception as e:
179                 self.logger.error(f"Canary deployment failed: {str(e)}")
180                 deployment_log["final_status"] = "error"
181                 deployment_log["error"] = str(e)
182
183             # Attempt cleanup

```

```

184         try:
185             await self._cleanup_failed_canary(canary_service_name)
186         except:
187             pass
188
189         raise
190
191     async def get_deployment_metrics(self, service_name: str) -> Optional[
192         DeploymentMetrics]:
193         """
194         Retrieve performance metrics for a deployment from Prometheus
195         """
196         try:
197             queries = {
198                 "success_rate": f'rate(kserve_request_total{{service_name="{
199                 service_name}",code!~"5.."}}[5m]) / rate(kserve_request_total{{service_name
200                ="{service_name}"}}[5m])',
201                 "avg_latency": f'histogram_quantile(0.50, rate(
202                 kserve_request_duration_seconds_bucket{{service_name="{service_name}"}}[5m]
203                 ) * 1000',
204                 "error_rate": f'rate(kserve_request_total{{service_name="{
205                 service_name}",code=~"5.."}}[5m]) / rate(kserve_request_total{{service_name
206                ="{service_name}"}}[5m])',
207                 "request_count": f'rate(kserve_request_total{{service_name="{
208                 service_name}"}}[5m]) * 300' # 5 minute rate * 300 seconds
209             }
210
211             metrics = {}
212             async with aiohttp.ClientSession() as session:
213                 for metric_name, query in queries.items():
214                     url = f"{self.prometheus_url}/api/v1/query"
215                     params = {"query": query}
216
217                     async with session.get(url, params=params) as response:
218                         if response.status == 200:
219                             data = await response.json()
220                             result = data.get("data", {}).get("result", [])
221
222                             if result:
223                                 value = float(result[0]["value"][1])
224                                 metrics[metric_name] = value
225                             else:
226                                 metrics[metric_name] = 0.0
227                         else:
228                             self.logger.warning(f"Failed to query {metric_name}
229                             ): {response.status}")
230                             metrics[metric_name] = 0.0
231
232             return DeploymentMetrics(
233                 success_rate=metrics.get("success_rate", 0.0),
234                 avg_latency_ms=metrics.get("avg_latency", 0.0),
235                 error_rate=metrics.get("error_rate", 0.0),
236                 request_count=int(metrics.get("request_count", 0)),
237                 timestamp=datetime.utcnow()
238             )
239
240         except Exception as e:
241             self.logger.error(f"Failed to retrieve metrics for {service_name}:
242             {str(e)}")
243             return None
244
245     def _evaluate_canary_performance(self,

```

```

236         baseline_metrics: Optional[DeploymentMetrics]
237     ],
238         canary_metrics: Optional[DeploymentMetrics],
239         final_evaluation: bool = False) -> Dict[str,
240         any]:
241         """
242         Evaluate canary performance against baseline and thresholds
243         """
244         if not canary_metrics:
245             return {
246                 "action": "rollback",
247                 "reason": "No canary metrics available",
248                 "confidence": 1.0
249             }
250
251         # Check absolute thresholds
252         threshold_checks = {
253             "success_rate": canary_metrics.success_rate >= self.config.
254             success_rate_threshold,
255             "latency": canary_metrics.avg_latency_ms <= self.config.
256             latency_threshold_ms,
257             "error_rate": canary_metrics.error_rate <= self.config.
258             error_rate_threshold,
259             "min_requests": canary_metrics.request_count >= 10 # Minimum
260             sample size
261         }
262
263         failed_checks = [check for check, passed in threshold_checks.items() if
264             not passed]
265
266         if failed_checks:
267             return {
268                 "action": "rollback",
269                 "reason": f"Failed threshold checks: {failed_checks}",
270                 "canary_metrics": canary_metrics.__dict__,
271                 "failed_checks": failed_checks,
272                 "confidence": 1.0
273             }
274
275         # Compare with baseline if available
276         if baseline_metrics and baseline_metrics.request_count >= 10:
277             relative_checks = {
278                 "success_rate_degradation": (canary_metrics.success_rate /
279                 baseline_metrics.success_rate) >= 0.99,
280                 "latency_increase": (canary_metrics.avg_latency_ms /
281                 baseline_metrics.avg_latency_ms) <= 1.2,
282                 "error_rate_increase": canary_metrics.error_rate <= (
283                 baseline_metrics.error_rate * 2.0 + 0.001)
284             }
285
286             failed_relative_checks = [check for check, passed in
287             relative_checks.items() if not passed]
288
289             if failed_relative_checks:
290                 # Calculate confidence based on sample size and magnitude of
291                 degradation
292                 confidence = min(1.0, canary_metrics.request_count / 100)
293
294                 return {
295                     "action": "rollback" if confidence > 0.7 else "hold",
296                     "reason": f"Performance degradation detected: {
297                     failed_relative_checks}",
298                     "canary_metrics": canary_metrics.__dict__,

```



```

286         "baseline_metrics": baseline_metrics.__dict__,
287         "failed_checks": failed_relative_checks,
288         "confidence": confidence
289     }
290
291     # Determine action based on evaluation type
292     if final_evaluation:
293         return {
294             "action": "promote",
295             "reason": "All performance checks passed",
296             "canary_metrics": canary_metrics.__dict__,
297             "confidence": 1.0
298         }
299     else:
300         return {
301             "action": "continue",
302             "reason": "Performance within acceptable range",
303             "canary_metrics": canary_metrics.__dict__,
304             "confidence": 0.8
305         }
306
307     async def _update_traffic_split(self, canary_service_name: str,
308                                     traffic_percent: int):
309         """Update traffic split for canary deployment"""
310         await self.kserve_integration.update_traffic_split(
311             canary_service_name,
312             {"canary": traffic_percent}
313         )
314
315     async def _rollback_canary(self, service_name: str, canary_service_name:
316                               str):
317         """Rollback canary deployment"""
318         try:
319             # Set canary traffic to 0
320             await self._update_traffic_split(canary_service_name, 0)
321
322             # Delete canary service after grace period
323             await asyncio.sleep(30)
324
325             self.kserve_integration.custom_client.
326             delete_namespaced_custom_object(
327                 group="serving.kserve.io",
328                 version="v1beta1",
329                 namespace=self.kserve_integration.namespace,
330                 plural="inferenceservices",
331                 name=canary_service_name
332             )
333
334             self.logger.info(f"Rolled back canary deployment: {
335                               canary_service_name}")
336
337         except Exception as e:
338             self.logger.error(f"Error during rollback: {str(e)}")
339
340     async def _promote_canary(self, service_name: str, canary_service_name: str
341                               , new_version: str):
342         """Promote canary to production"""
343         try:
344             # Get canary service configuration
345             canary_service = self.kserve_integration.custom_client.
346             get_namespaced_custom_object(
347                 group="serving.kserve.io",
348                 version="v1beta1",

```

```

343         namespace=self.kserve_integration.namespace,
344         plural="inferenceservices",
345         name=canary_service_name
346     )
347
348     # Update production service with canary configuration
349     patch_body = {
350         "spec": canary_service["spec"],
351         "metadata": {
352             "labels": canary_service["metadata"]["labels"],
353             "annotations": {
354                 **canary_service["metadata"]["annotations"],
355                 "promotion.timestamp": datetime.utcnow().isoformat(),
356                 "promoted.from": canary_service_name
357             }
358         }
359     }
360
361     # Remove canary-specific configurations
362     if "canaryTrafficPercent" in patch_body["spec"]["predictor"]:
363         del patch_body["spec"]["predictor"]["canaryTrafficPercent"]
364
365     # Update production service
366     self.kserve_integration.custom_client.
patch_namespaced_custom_object(
367         group="serving.kserve.io",
368         version="v1beta1",
369         namespace=self.kserve_integration.namespace,
370         plural="inferenceservices",
371         name=service_name,
372         body=patch_body
373     )
374
375     # Clean up canary service
376     await asyncio.sleep(60) # Wait for traffic to shift
377     await self._cleanup_failed_canary(canary_service_name)
378
379     self.logger.info(f"Successfully promoted {canary_service_name} to
production")
380
381     except Exception as e:
382         self.logger.error(f"Error during promotion: {str(e)}")
383         raise
384
385     async def _cleanup_failed_canary(self, canary_service_name: str):
386         """Clean up failed canary deployment"""
387         try:
388             self.kserve_integration.custom_client.
delete_namespaced_custom_object(
389                 group="serving.kserve.io",
390                 version="v1beta1",
391                 namespace=self.kserve_integration.namespace,
392                 plural="inferenceservices",
393                 name=canary_service_name
394             )
395             self.logger.info(f"Cleaned up canary service: {canary_service_name}
")
396 \subsubsection{Blue-Green Deployments}
397
398 Implement zero-downtime deployments with instant rollback capabilities:
399
400 \begin{lstlisting}[language=python, caption=Blue-Green Deployment
Implementation]

```

```

401 class BlueGreenDeploymentManager:
402     """
403     Manages blue-green deployments for zero-downtime model updates
404     """
405
406     def __init__(self, kserve_integration: MLflowKServeIntegration):
407         self.kserve_integration = kserve_integration
408         self.logger = logging.getLogger(__name__)
409
410     async def deploy_blue_green(self,
411                                model_name: str,
412                                new_version: str,
413                                service_name: str,
414                                validation_tests: List[Callable] = None) -> Dict[
415                                str, Any]:
416         """
417         Execute blue-green deployment with automated validation
418         """
419
420         deployment_result = {
421             "model_name": model_name,
422             "new_version": new_version,
423             "service_name": service_name,
424             "start_time": datetime.utcnow(),
425             "status": "in_progress"
426         }
427
428         try:
429             # Step 1: Deploy green environment
430             green_service_name = f"{service_name}-green"
431
432             self.logger.info(f"Deploying green environment: {green_service_name}")
433
434             green_deployment = self.kserve_integration.create_inference_service(
435                 (
436                     model_name=model_name,
437                     model_version=new_version,
438                     service_name=green_service_name
439                 )
440             )
441
442             # Step 2: Wait for green environment to be ready
443             await self._wait_for_service_ready(green_service_name)
444
445             # Step 3: Run validation tests against green environment
446             if validation_tests:
447                 validation_results = await self._run_validation_tests(
448                     green_service_name,
449                     validation_tests
450                 )
451
452                 if not validation_results["all_passed"]:
453                     # Cleanup and abort
454                     await self._cleanup_service(green_service_name)
455                     deployment_result["status"] = "validation_failed"
456                     deployment_result["validation_results"] =
457                     validation_results
458
459                     return deployment_result
460
461             # Step 4: Switch traffic from blue to green
462             self.logger.info("Switching traffic from blue to green")
463
464             # Get current blue service configuration

```

```

460         blue_service = self.kserve_integration.custom_client.
get_namespaced_custom_object(
461             group="serving.kserve.io",
462             version="v1beta1",
463             namespace=self.kserve_integration.namespace,
464             plural="inferenceservices",
465             name=service_name
466         )
467
468         # Backup blue configuration
469         blue_backup_name = f"{service_name}-blue-backup-{int(datetime.
utcnow().timestamp())}"
470         deployment_result["blue_backup_name"] = blue_backup_name
471
472         # Rename current service to backup
473         await self._rename_service(service_name, blue_backup_name)
474
475         # Rename green service to production
476         await self._rename_service(green_service_name, service_name)
477
478         # Step 5: Monitor post-deployment metrics
479         post_deployment_metrics = await self._monitor_post_deployment(
service_name)
480         deployment_result["post_deployment_metrics"] =
post_deployment_metrics
481
482         # Step 6: Cleanup old blue environment after successful deployment
483         await asyncio.sleep(300) # Wait 5 minutes before cleanup
484         await self._cleanup_service(blue_backup_name)
485
486         deployment_result["status"] = "completed"
487         deployment_result["end_time"] = datetime.utcnow()
488
489         self.logger.info(f"Blue-green deployment completed successfully for
{model_name}")
490
491         return deployment_result
492
493     except Exception as e:
494         self.logger.error(f"Blue-green deployment failed: {str(e)}")
495         deployment_result["status"] = "failed"
496         deployment_result["error"] = str(e)
497
498         # Attempt rollback if switch was attempted
499         if "blue_backup_name" in deployment_result:
500             try:
501                 await self.rollback_blue_green(service_name,
deployment_result["blue_backup_name"])
502                 deployment_result["rollback_attempted"] = True
503             except:
504                 deployment_result["rollback_failed"] = True
505
506             raise
507
508     async def rollback_blue_green(self, service_name: str, blue_backup_name:
str):
509         """
510         Rollback blue-green deployment to previous version
511         """
512         try:
513             self.logger.info(f"Rolling back blue-green deployment for {
service_name}")
514

```

```

515         # Delete current green service
516         await self._cleanup_service(service_name)
517
518         # Restore blue service
519         await self._rename_service(blue_backup_name, service_name)
520
521         self.logger.info("Blue-green rollback completed successfully")
522
523     except Exception as e:
524         self.logger.error(f"Blue-green rollback failed: {str(e)}")
525         raise
526
527     async def _wait_for_service_ready(self, service_name: str, timeout_seconds:
528         int = 300):
529         """Wait for service to be ready"""
530         start_time = datetime.utcnow()
531
532         while (datetime.utcnow() - start_time).total_seconds() <
533             timeout_seconds:
534             try:
535                 service = self.kserve_integration.custom_client.
536                 get_namespaced_custom_object(
537                     group="serving.kserve.io",
538                     version="v1beta1",
539                     namespace=self.kserve_integration.namespace,
540                     plural="inferenceservices",
541                     name=service_name
542                 )
543
544                 # Check if service is ready
545                 conditions = service.get("status", {}).get("conditions", [])
546                 ready_condition = next((c for c in conditions if c["type"] == "
547                     Ready"), None)
548
549                 if ready_condition and ready_condition["status"] == "True":
550                     self.logger.info(f"Service {service_name} is ready")
551                     return
552
553             except Exception as e:
554                 self.logger.debug(f"Waiting for service readiness: {str(e)}")
555
556             await asyncio.sleep(10)
557
558         raise TimeoutError(f"Service {service_name} did not become ready within
559             {timeout_seconds} seconds")
560
561     async def _run_validation_tests(self, service_name: str, validation_tests:
562         List[callable]) -> Dict[str, Any]:
563         """Run validation tests against deployed service"""
564         results = {
565             "all_passed": True,
566             "test_results": [],
567             "execution_time": datetime.utcnow()
568         }
569
570         for test_func in validation_tests:
571             try:
572                 test_result = await test_func(service_name)
573                 results["test_results"].append({
574                     "test_name": test_func.__name__,
575                     "status": "passed" if test_result else "failed",
576                     "result": test_result
577                 })

```

```

572         if not test_result:
573             results["all_passed"] = False
574
575     except Exception as e:
576         results["test_results"].append({
577             "test_name": test_func.__name__,
578             "status": "error",
579             "error": str(e)
580         })
581         results["all_passed"] = False
582
583     return results
584
585 \subsection{A/B Testing Framework}
586
587 Implement statistical A/B testing for model variants:
588
589 \begin{lstlisting}[language=python, caption=A/B Testing Implementation]
590 import scipy.stats as stats
591 from typing import Tuple
592 import pandas as pd
593
594 class ABTestingFramework:
595     """
596     Statistical A/B testing framework for model variants
597     """
598
599     def __init__(self,
600                 kserve_integration: MLflowKServeIntegration,
601                 prometheus_url: str,
602                 significance_level: float = 0.05,
603                 minimum_sample_size: int = 1000,
604                 test_duration_hours: int = 72):
605
606         self.kserve_integration = kserve_integration
607         self.prometheus_url = prometheus_url
608         self.significance_level = significance_level
609         self.minimum_sample_size = minimum_sample_size
610         self.test_duration_hours = test_duration_hours
611         self.logger = logging.getLogger(__name__)
612
613     async def setup_ab_test(self,
614                             control_model: Dict[str, str],
615                             variant_model: Dict[str, str],
616                             traffic_split: int = 50,
617                             success_metric: str = "conversion_rate") -> Dict[str,
618                                     Any]:
619         """
620         Setup A/B test between two model variants
621         """
622
623         test_config = {
624             "test_id": f"ab_test_{int(datetime.utcnow().timestamp())}",
625             "control_model": control_model,
626             "variant_model": variant_model,
627             "traffic_split": traffic_split,
628             "success_metric": success_metric,
629             "start_time": datetime.utcnow(),
630             "status": "active"
631         }
632
633         try:

```

```

634         # Deploy control version (if not already deployed)
635         control_service_name = f"{control_model['name']}-control-{{
test_config['test_id']}}"
636
637         if not await self._service_exists(control_service_name):
638             await self.kserve_integration.create_inference_service(
639                 model_name=control_model['name'],
640                 model_version=control_model['version'],
641                 service_name=control_service_name
642             )
643
644         # Deploy variant version
645         variant_service_name = f"{variant_model['name']}-variant-{{
test_config['test_id']}}"
646
647         await self.kserve_integration.create_inference_service(
648             model_name=variant_model['name'],
649             model_version=variant_model['version'],
650             service_name=variant_service_name
651         )
652
653         # Configure traffic splitting
654         await self._configure_ab_traffic(
655             control_service_name,
656             variant_service_name,
657             traffic_split
658         )
659
660         test_config["control_service"] = control_service_name
661         test_config["variant_service"] = variant_service_name
662
663         self.logger.info(f"A/B test setup completed: {{test_config['test_id
',]}}")
664
665         return test_config
666
667     except Exception as e:
668         self.logger.error(f"Failed to setup A/B test: {{str(e)}}")
669         test_config["status"] = "failed"
670         test_config["error"] = str(e)
671         raise
672
673     async def analyze_ab_test(self, test_config: Dict[str, Any]) -> Dict[str,
Any]:
674         """
675         Analyze A/B test results and make statistical conclusions
676         """
677         try:
678             # Collect metrics for both variants
679             control_metrics = await self._collect_ab_metrics(
680                 test_config["control_service"],
681                 test_config["start_time"]
682             )
683
684             variant_metrics = await self._collect_ab_metrics(
685                 test_config["variant_service"],
686                 test_config["start_time"]
687             )
688
689             # Perform statistical analysis
690             analysis_result = self._perform_statistical_analysis(
691                 control_metrics,
692                 variant_metrics,

```

```

693         test_config["success_metric"]
694     )
695
696     # Generate recommendations
697     recommendation = self._generate_recommendation(analysis_result,
test_config)
698
699     result = {
700         "test_id": test_config["test_id"],
701         "analysis_timestamp": datetime.utcnow(),
702         "control_metrics": control_metrics,
703         "variant_metrics": variant_metrics,
704         "statistical_analysis": analysis_result,
705         "recommendation": recommendation,
706         "test_duration_hours": (datetime.utcnow() - test_config["
start_time"])).total_seconds() / 3600
707     }
708
709     return result
710
711 except Exception as e:
712     self.logger.error(f"Failed to analyze A/B test: {str(e)}")
713     raise
714
715 def _perform_statistical_analysis(self,
716                                 control_metrics: Dict,
717                                 variant_metrics: Dict,
718                                 success_metric: str) -> Dict[str, Any]:
719     """
720     Perform statistical significance testing
721     """
722
723     # Extract success counts and total counts
724     control_successes = control_metrics.get(f"{success_metric}_count", 0)
725     control_total = control_metrics.get("total_requests", 0)
726     variant_successes = variant_metrics.get(f"{success_metric}_count", 0)
727     variant_total = variant_metrics.get("total_requests", 0)
728
729     if control_total == 0 or variant_total == 0:
730         return {
731             "test_type": "insufficient_data",
732             "statistical_significance": False,
733             "p_value": None,
734             "confidence_interval": None,
735             "effect_size": None
736         }
737
738     # Calculate conversion rates
739     control_rate = control_successes / control_total
740     variant_rate = variant_successes / variant_total
741
742     # Perform two-proportion z-test
743     count = np.array([control_successes, variant_successes])
744     nobs = np.array([control_total, variant_total])
745
746     # Calculate z-statistic and p-value
747     z_stat, p_value = stats.proportions_ztest(count, nobs)
748
749     # Calculate confidence interval for difference
750     pooled_rate = (control_successes + variant_successes) / (control_total
+ variant_total)
751     se_diff = np.sqrt(pooled_rate * (1 - pooled_rate) * (1/control_total +
1/variant_total))

```



```

752     rate_diff = variant_rate - control_rate
753
754     margin_of_error = stats.norm.ppf(1 - self.significance_level/2) *
se_diff
755     ci_lower = rate_diff - margin_of_error
756     ci_upper = rate_diff + margin_of_error
757
758     # Calculate effect size (Cohen's h)
759     effect_size = 2 * (np.arcsin(np.sqrt(variant_rate)) - np.arcsin(np.sqrt
(control_rate)))
760
761     return {
762         "test_type": "two_proportion_z_test",
763         "control_rate": control_rate,
764         "variant_rate": variant_rate,
765         "rate_difference": rate_diff,
766         "relative_improvement": (rate_diff / control_rate) * 100 if
control_rate > 0 else 0,
767         "z_statistic": z_stat,
768         "p_value": p_value,
769         "statistical_significance": p_value < self.significance_level,
770         "confidence_interval": (ci_lower, ci_upper),
771         "effect_size": effect_size,
772         "sample_sizes": {"control": control_total, "variant": variant_total
}
773     }
774
775     def _generate_recommendation(self,
776                                analysis_result: Dict[str, Any],
777                                test_config: Dict[str, Any]) -> Dict[str, Any]:
778         """
779         Generate actionable recommendations based on test results
780         """
781
782         if analysis_result["test_type"] == "insufficient_data":
783             return {
784                 "action": "continue_test",
785                 "reason": "Insufficient data for statistical analysis",
786                 "required_sample_size": self.minimum_sample_size
787             }
788
789         # Check minimum sample size requirement
790         min_sample_met = all(
791             size >= self.minimum_sample_size
792             for size in analysis_result["sample_sizes"].values()
793         )
794
795         if not min_sample_met:
796             return {
797                 "action": "continue_test",
798                 "reason": "Minimum sample size not reached",
799                 "current_samples": analysis_result["sample_sizes"],
800                 "required_sample_size": self.minimum_sample_size
801             }
802
803         # Check test duration
804         test_duration = (datetime.utcnow() - test_config["start_time"]).
total_seconds() / 3600
805         if test_duration < self.test_duration_hours:
806             return {
807                 "action": "continue_test",
808                 "reason": f"Test duration ({test_duration:.1f}h) below minimum
({self.test_duration_hours}h)",

```

```

809         "current_duration_hours": test_duration,
810         "required_duration_hours": self.test_duration_hours
811     }
812
813     # Make recommendation based on statistical results
814     if analysis_result["statistical_significance"]:
815         if analysis_result["rate_difference"] > 0:
816             return {
817                 "action": "deploy_variant",
818                 "reason": "Variant shows statistically significant
improvement",
819                 "improvement": f"{analysis_result['relative_improvement
']:.2f}%",
820                 "confidence": f"{{(1 - self.significance_level) * 100:.0f}}%"
821             }
822         else:
823             return {
824                 "action": "keep_control",
825                 "reason": "Control performs significantly better than
variant",
826                 "degradation": f"{analysis_result['relative_improvement
']:.2f}%",
827                 "confidence": f"{{(1 - self.significance_level) * 100:.0f}}%"
828             }
829         else:
830             return {
831                 "action": "no_significant_difference",
832                 "reason": "No statistically significant difference detected",
833                 "p_value": analysis_result["p_value"],
834                 "recommendation": "Consider other factors like cost, complexity
, or business requirements"
835             }
836
837 # Example usage for A/B testing
838 async def example_ab_test():
839     """Example A/B test implementation"""
840
841     # Initialize framework
842     ab_framework = ABTestingFramework(
843         kserve_integration=integration,
844         prometheus_url="http://prometheus.monitoring.svc.cluster.local:9090"
845     )
846
847     # Setup A/B test
848     test_config = await ab_framework.setup_ab_test(
849         control_model={"name": "recommendation-model", "version": "1.2.0"},
850         variant_model={"name": "recommendation-model", "version": "1.3.0"},
851         traffic_split=50,
852         success_metric="click_through_rate"
853     )
854
855     # Monitor test for specified duration
856     while True:
857         analysis = await ab_framework.analyze_ab_test(test_config)
858
859         if analysis["recommendation"]["action"] != "continue_test":
860             print(f"Test completed with recommendation: {analysis['
recommendation']['action']}")
861             break
862
863             print(f"Test continuing... Current improvement: {analysis['
statistical_analysis'].get('relative_improvement', 0):.2f}%")

```

```
864 await asyncio.sleep(3600) # Check every hour
```

Listing 13: Automated Canary Deployment Controller

This comprehensive model deployment and serving section provides production-ready implementations for advanced deployment patterns including canary deployments, blue-green deployments, and statistical A/B testing frameworks, all integrated with KServe and MLflow for complete MLOps workflow management.

4 Monitoring and Observability

4.1 Model Performance Monitoring

Implementing comprehensive monitoring is crucial for maintaining model performance in production:

```
1 import prometheus_client
2 from prometheus_client import Counter, Histogram, Gauge
3 import numpy as np
4 from scipy import stats
5
6 class ModelMonitor:
7     def __init__(self):
8         # Prometheus metrics
9         self.prediction_counter = Counter(
10             'model_predictions_total',
11             'Total number of predictions made',
12             ['model_name', 'version']
13         )
14
15         self.prediction_latency = Histogram(
16             'model_prediction_duration_seconds',
17             'Model prediction latency',
18             ['model_name', 'version']
19         )
20
21         self.model_accuracy = Gauge(
22             'model_accuracy_score',
23             'Current model accuracy',
24             ['model_name', 'version']
25         )
26
27         self.drift_score = Gauge(
28             'model_drift_score',
29             'Data drift detection score',
30             ['model_name', 'feature']
31         )
32
33     def log_prediction(self, model_name, version, latency):
34         """Log prediction metrics"""
35         self.prediction_counter.labels(
36             model_name=model_name,
37             version=version
38         ).inc()
39
40         self.prediction_latency.labels(
41             model_name=model_name,
42             version=version
43         ).observe(latency)
44
45     def detect_data_drift(self, reference_data, current_data, feature_name):
46         """Detect data drift using Kolmogorov-Smirnov test"""
```

```

47         try:
48             # Perform KS test
49             ks_statistic, p_value = stats.ks_2samp(reference_data, current_data
50         )
51
52             # Update drift metric
53             self.drift_score.labels(
54                 model_name="production_model",
55                 feature=feature_name
56             ).set(ks_statistic)
57
58             # Alert if significant drift detected
59             if p_value < 0.05:
60                 self.send_drift_alert(feature_name, ks_statistic, p_value)
61
62             return ks_statistic, p_value
63
64         except Exception as e:
65             print(f"Error detecting drift for {feature_name}: {str(e)}")
66             return None, None
67
68     def send_drift_alert(self, feature_name, ks_statistic, p_value):
69         """Send alert when data drift is detected"""
70         alert_message = f"""
71         Data Drift Alert!
72         Feature: {feature_name}
73         KS Statistic: {ks_statistic:.4f}
74         P-value: {p_value:.4f}
75         Recommendation: Review model performance and consider retraining
76         """
77         # Integration with alerting system (Slack, PagerDuty, etc.)
78         print(alert_message)

```

Listing 14: Model Monitoring Implementation

5 Best Practices and Recommendations

5.1 Pipeline Design Principles

1. **Modularity:** Design pipeline components as independent, reusable modules that can be easily tested and maintained.
2. **Reproducibility:** Ensure all experiments and deployments are fully reproducible through proper versioning of code, data, and dependencies.
3. **Scalability:** Design pipelines to handle varying workloads and data volumes without manual intervention.
4. **Monitoring:** Implement comprehensive monitoring at every stage of the pipeline to quickly identify and resolve issues.
5. **Security:** Apply security best practices including proper authentication, authorization, and data encryption.

5.2 Performance Optimization

- Use GPU acceleration for training intensive models

- Implement efficient data loading and preprocessing
- Optimize model serving with batching and caching strategies
- Monitor resource utilization and scale components as needed
- Implement model quantization and pruning for inference optimization

6 Troubleshooting Common Issues

6.1 Pipeline Failures

Common issues and their solutions:

| Issue | Symptoms | Solution |
|----------------------|--------------------------------------|---|
| Resource limitations | Pipeline steps timing out or failing | Increase resource requests/limits |
| Data access issues | Permission denied errors | Check RBAC and storage permissions |
| Model convergence | Poor model performance | Adjust hyperparameters and data quality |
| Deployment failures | Service unavailable errors | Verify KServe configuration |

Table 1: Common Pipeline Issues and Solutions

7 Conclusion

This comprehensive guide demonstrates how to build production-ready MLOps pipelines using Kubeflow and MLflow. The combination of these technologies provides a robust foundation for managing the complete machine learning lifecycle, from experimentation to production deployment.

Key benefits of this approach include:

- **Standardization:** Consistent workflows across teams and projects
- **Scalability:** Kubernetes-native scaling capabilities
- **Reproducibility:** Complete experiment and deployment tracking
- **Monitoring:** Comprehensive observability and alerting
- **Collaboration:** Shared infrastructure and knowledge base

As MLOps practices continue to evolve, this foundation provides the flexibility to adapt and integrate new tools and methodologies while maintaining operational excellence.

8 References

1. Kubeflow Documentation. <https://www.kubeflow.org/docs/>
2. MLflow Documentation. <https://mlflow.org/docs/latest/index.html>
3. KServe Documentation. <https://kserve.github.io/website/>

4. Kubernetes Documentation. <https://kubernetes.io/docs/>
5. Sculley, D., et al. "Hidden Technical Debt in Machine Learning Systems." NIPS 2015.

9 About the Author

Vladimir Ovcharov is an MLOps Engineer and ML Systems Architect with 8+ years of experience in building production-ready machine learning infrastructure. He specializes in designing scalable ML pipelines, implementing robust monitoring systems, and optimizing model deployment workflows.

Contact: vladimir.ovcharov@example.com